

000  
001  
002  
003  
004  
005  
006  
007  
008  
009  
010  
011  
012  
013  
014  
015  
016  
017  
018  
019  
020  
021  
022  
023  
024  
025  
026  
027  
028  
029  
030  
031  
032  
033  
034  
035  
036  
037  
038  
039  
040  
041  
042  
043  
044  
045  
046  
047  
048  
049  
050  
051  
052  
053

THE COOPER UNION  
ALBERT NERKEN SCHOOL OF ENGINEERING

A Deep Partitioned Autoencoder  
for De-Noising Live Audio

by  
Ethan Lusterman

A thesis submitted in partial fulfillment  
of the requirements for the degree of  
Master of Engineering

September 2016

Professor Sam Keene, Advisor

054 THE COOPER UNION FOR THE  
055  
056 ADVANCEMENT OF SCIENCE AND ART  
057  
058  
059  
060

061 ALBERT NERKEN SCHOOL OF ENGINEERING  
062  
063  
064  
065  
066  
067  
068  
069

070 This thesis was prepared under the direction of the Can-  
071 didate's Thesis Advisor and has received approval. It was  
072 submitted to the Dean of the School of Engineering and  
073 the full Faculty, and was approved as partial fulfillment of  
074 the requirements for the degree of Master of Engineering.  
075  
076  
077  
078  
079  
080  
081  
082  
083  
084  
085  
086

087 \_\_\_\_\_  
088 Dean, School of Engineering                      Date  
089  
090  
091  
092  
093  
094

095 \_\_\_\_\_  
096 Prof. Sam Keene, Thesis Advisor                      Date  
097  
098  
099  
100  
101  
102  
103  
104  
105  
106  
107

108  
109  
110  
111  
112  
113  
114  
115  
116  
117  
118  
119  
120  
121  
122  
123  
124  
125  
126  
127  
128  
129  
130  
131  
132  
133  
134  
135  
136  
137  
138  
139  
140  
141  
142  
143  
144  
145  
146  
147  
148  
149  
150  
151  
152  
153  
154  
155  
156  
157  
158  
159  
160  
161

# Acknowledgements

Ack example

162  
163  
164  
165  
166  
167  
168  
169  
170  
171  
172  
173  
174  
175  
176  
177  
178  
179  
180  
181  
182  
183  
184  
185  
186  
187  
188  
189  
190  
191  
192  
193  
194  
195  
196  
197  
198  
199  
200  
201  
202  
203  
204  
205  
206  
207  
208  
209  
210  
211  
212  
213  
214  
215

# Abstract

## Abstract

216  
217  
218  
219  
220  
221  
222  
223  
224  
225  
226  
227  
228  
229  
230  
231  
232  
233  
234  
235  
236  
237  
238  
239  
240  
241  
242  
243  
244  
245  
246  
247  
248  
249  
250  
251  
252  
253  
254  
255  
256  
257  
258  
259  
260  
261  
262  
263  
264  
265  
266  
267  
268  
269

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Background</b>	<b>2</b>
2.1	Machine Learning . . . . .	2
2.1.1	Regression . . . . .	3
2.1.2	Overfitting and Curse of Dimensionality . . . . .	3
2.1.3	Loss functions and Regularization . . . . .	3
2.1.4	Gradient Stuff? . . . . .	3
2.2	Neural Networks . . . . .	3
2.2.1	Dense Layer . . . . .	3
2.2.2	Convolutional Layer . . . . .	3
2.2.3	Nonlinearity Choice . . . . .	3
2.3	Signals and Systems . . . . .	3
2.3.1	Signals . . . . .	3
2.3.2	Convolution . . . . .	4
2.3.3	Frequency Transforms . . . . .	5
2.3.4	Windowing and Perfect Reconstruction . . . . .	6
2.3.5	Noise and Signal-to-Noise Ratio . . . . .	7
2.3.6	Magnitude and Phase Spectrum . . . . .	7
<b>3</b>	<b>System Description</b>	<b>8</b>
<b>4</b>	<b>Results</b>	<b>9</b>
<b>5</b>	<b>Conclusions and Future Work</b>	<b>10</b>
5.1	Conclusions . . . . .	10
5.2	Future Work . . . . .	10
5.2.1	Models . . . . .	10
5.2.2	Data . . . . .	10

270  
271  
272  
273  
274  
275  
276  
277  
278  
279  
280  
281  
282  
283  
284  
285  
286  
287  
288  
289  
290  
291  
292  
293  
294  
295  
296  
297  
298  
299  
300  
301  
302  
303  
304  
305  
306  
307  
308  
309  
310  
311  
312  
313  
314  
315  
316  
317  
318  
319  
320  
321  
322  
323

## List of Figures

324	Table of Nomenclature
325	
326	
327	
328	
329	
330	
331	
332	
333	
334	
335	
336	
337	
338	
339	
340	
341	
342	
343	
344	
345	
346	
347	
348	
349	
350	
351	
352	
353	
354	
355	
356	
357	
358	
359	
360	
361	
362	
363	
364	
365	
366	
367	
368	
369	
370	
371	
372	
373	
374	
375	
376	
377	

# 1 Introduction

Advances in smartphone technology have led to smaller devices with more powerful audio hardware, allowing for common consumers to make higher quality recordings. However, recorded speech and music are subject to noisy conditions, often hampering intelligibility and listenability. The goal of denoising audio recordings is to improve intelligibility and perceived quality. A variety of applications of audio denoising exist, including listening to a recording of a band or an artist’s live performance in a noisy crowd, or listening to a recorded conversation or speech under noisy conditions.

A common technique for denoising involves the use of deep neural networks (DNN). [PARIS] Advances in parallel graphics processing units (GPU) and in machine learning algorithms have allowed for training deeper networks faster, utilizing more hidden layers with more neurons.

Prior work in denoising audio has involved access to noise-free training data. Since common consumers do not often have access to clean audio, we seek to denoise without the use of clean audio.

In this thesis, we compare several neural network architectures and problem scenarios, ranging from data input types, level of noise, depth of network, training objectives, and more. In Chapter 2, we present background information on machine learning, neural networks, and signal processing as well as prior work in audio denoising. In Chapter 3, we detail all considered network architectures. In Chapter 4, we compare results from different data inputs, levels of noise, network architectures, and training objectives and discuss methods of evaluation. Finally, we make conclusions and recommendations for future work in Chapter 5.



## 2 Background

### 2.1 Machine Learning

Machine learning involves the use of computer algorithms to make decisions based on training data. Generally, this falls into categorizing input data (classification) or determining a mathematical function to determine a continuous output given an input (regression). Popular classification examples include recognizing handwritten digits (MNIST) as well as determining whether an image contains a cat or a dog. (REF) An example of a regression problem is determining the temperature given a set of input features (humidity, latitude, longitude, date, etc.).

Problems where training data contain input data vectors as well as the correct output vectors (targets) are known as supervised learning problems. Training a model to denoise audio where noise was introduced to the clean audio would be a supervised learning problem. On the other hand, training a model to denoise audio where the underlying clean signal is not known is an unsupervised learning problem. Different loss functions and neural network architectures can be exploited to accomplish denoising without the clean data.

For the purposes of this thesis, we use machine learning to determine an underlying nonlinear function that removes noise from time slices of audio (i.e. regression). These slices can then be pieced back together through overlap-add resynthesis. To clarify, this is a general linear model that maps an input noisy audio vector  $y[n] = x[n] + N[n]$  to  $\tilde{x}[n]$ , a target denoised audio vector, where  $x[n]$  is the underlying clean signal and  $N[n]$  is the additive background noise.

### 2.1.1 Regression

A classical regression technique is linear regression, where one or more independent variables  $x_i$  are used to determine a scalar dependent variable  $y$ . The case of a single independent variable  $x$  is known as simple linear regression. On the other hand, the case of estimating A canonical example would be estimating a sine wave  $x[n]$  given noisy samples

### 2.1.2 Overfitting and Curse of Dimensionality

### 2.1.3 Loss functions and Regularization

### 2.1.4 Gradient Stuff?

## 2.2 Neural Networks

### 2.2.1 Dense Layer

### 2.2.2 Convolutional Layer

### 2.2.3 Nonlinearity Choice

## 2.3 Signals and Systems

Domain knowledge of discrete audio signals and systems better informs our decisions for an audio denoising system, so some background information on signals and systems as it pertains to this thesis is detailed below.

### 2.3.1 Signals

We deal exclusively with discrete-time audio signals in this thesis. A discrete-time audio signal  $x[n]$  is represented as a sequence of numbers (samples), where each integer-valued slot  $n$  in the sequence corresponds to a unit of time based on the sampling frequency  $f_s$ . This comes from sampling the continuous-time audio signal  $x_c(t)$ :

$$x[n] = x_c(nT) \quad (1)$$

where  $T = 1/f_s$ . For example, a 1-second speech signal sampled at 8kHz has 8000 samples. Furthermore, digital signals also have discrete valued sample amplitudes. For the purposes of this thesis, the bit depths of computers we use for analysis are high enough to allow for perfect reconstruction between continuous-time signals and digital signals.

We also assume signals collected have been properly sampled according to the Nyquist-Shannon sampling theorem, which states that a discrete-time signal must be sampled at at least twice the highest frequency present in the signal to prevent aliasing of different frequencies. For example, speech signals generally have information up to 8kHz, so many speech signals are sampled at 16kHz. Music is more complex in that signals often span up to about 20kHz, so CD quality recordings are often sampled at 44.1kHz or higher. For this thesis, we use recordings sampled at 44.1kHz or lower.

### 2.3.2 Convolution

The discrete-time convolution operation takes two sequences  $x[n]$  and  $h[n]$  and outputs a third sequence  $y[n] = x[n] * h[n]$ :

$$y[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k] \quad (2)$$

Convolution is commutative, so  $x[n] * h[n] = h[n] * x[n]$  holds true.

A linear, time-invariant (LTI) system is characterized by its impulse response  $h[n]$ , which allows us to determine samples  $y[n]$  when  $x[n]$  is subject to  $h[n]$ . For the purposes of this thesis, our underlying clean signal  $x[n]$  might be

subject to the conditions of an acoustic environment  $h[n]$  and crowd noise  $N[n]$ :

$$y[n] = h[n] * x[n] + N[n] \quad (3)$$

In this scenario, our system would attempt to recover  $h[n] * x[n]$  and possibly even  $x[n]$  if the acoustic environment were deemed “noisy enough” due to echo and reverberation.

One of our proposed systems also incorporates convolutional neural networks (CNN) which use convolutions between frames of samples instead of simple linear combinations (discussed later).

### 2.3.3 Frequency Transforms

In some of our proposed systems, we use a frequency transformed version of the input signal as a preprocessing step to the system input. While no new information is gained from transforming the input, networks often respond better to determining the value of the magnitude of varying frequencies at a time slice instead of the individual time samples.

The frequency transform we use in this thesis is the discrete-time Fourier transform (DTFT). A sequence of  $N$  discrete-time samples is transformed into another sequence of  $N$  samples where each index then corresponds to a frequency bin. The DTFT  $X[k]$  of a signal  $x[n]$  is given by the following:

$$X[k] = \sum_{n=0}^{N-1} x[n] W_N^{kn} \quad (4)$$

where the twiddle factor  $W_N$  is given by  $W_N = e^{-j(2\pi/N)}$ . Then the reconstruction of  $x[n]$  from  $X[k]$  is given by:

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] W_N^{-kn} \quad (5)$$

In this thesis, we also exploit the main duality between the time and frequency domain using the convolution theorem, which states that convolution in time is equivalent to multiplication in frequency and vice versa:

$$\mathcal{F}\{h[n] * x[n]\} = H[k]X[k] \quad (6)$$

$$\mathcal{F}^{-1}\{H[k] * X[k]\} = h[n]x[n] \quad (7)$$

This allows us to effectively treat our network as a non-linear filter that can denoise small time/frequency slices of our noisy signal, which can then be pieced back together using overlap-add resynthesis. We detail this in the next section.

### 2.3.4 Windowing and Perfect Reconstruction

To window a signal is to multiply a window function  $w[n]$  by the frame, i.e.  $w[n]x[n]$  over the frame length  $N$ . Because we are training a network to denoise small segments of a larger audio signal, we window the signal segments. This accommodates the finite-length requirement of the DTFT and helps to prevent spectral leakage. [DSPBOOK]

Also, to be able to properly reconstruct our signal, we use a window function and corresponding overlapping frame percentage to accomplish perfect reconstruction. The corresponding overlapping frame percentage is set such that the window sums to a constant for all time. For example, a rectangular window  $w[n] = 1$  over an interval of length  $N$  has an overlap of 0% to sum to a constant 1 for all time. Another popular window is the Hanning window, defined over an interval  $N$  by the following:

$$w[n] = \frac{1}{2} \left( 1 - \cos \left( \frac{2\pi n}{N-1} \right) \right) \quad (8)$$

For the Hann window, the perfect reconstruction overlap is a frame length of  $N = 50\%$ .

### 2.3.5 Noise and Signal-to-Noise Ratio

Since we are trying to denoise audio signals, we must discuss how we measure noise. One of the most common measures of degradation of signal quality from additive noise is signal-to-noise ratio (SNR), defined as the ratio of signal variance to noise variance. [DSP] For the signal  $y[n] = x[n] + N[n]$ , where  $x[n]$  is the signal of interest and  $N[n]$  is the additive noise, the SNR is defined as

$$SNR = \frac{\sigma_x^2}{\sigma_n^2} \quad (9)$$

where  $\sigma^2$  refers to the variance of the signal in question over some time interval. For the purposes of this thesis, we achieve desired a desired SNR for a simulation by scaling the noise to match the variance to the signal, then scaling the noise or the signal to achieve the desired SNR.

### 2.3.6 Magnitude and Phase Spectrum

756  
757  
758  
759  
760  
761  
762  
763  
764  
765  
766  
767  
768  
769  
770  
771  
772  
773  
774  
775  
776  
777  
778  
779  
780  
781  
782  
783  
784  
785  
786  
787  
788  
789  
790  
791  
792  
793  
794  
795  
796  
797  
798  
799  
800  
801  
802  
803  
804  
805  
806  
807  
808  
809

### 3 System Description

810  
811  
812  
813  
814  
815  
816  
817  
818  
819  
820  
821  
822  
823  
824  
825  
826  
827  
828  
829  
830  
831  
832  
833  
834  
835  
836  
837  
838  
839  
840  
841  
842  
843  
844  
845  
846  
847  
848  
849  
850  
851  
852  
853  
854  
855  
856  
857  
858  
859  
860  
861  
862  
863

## 4 Results



864  
865  
866  
867  
868  
869  
870  
871  
872  
873  
874  
875  
876  
877  
878  
879  
880  
881  
882  
883  
884  
885  
886  
887  
888  
889  
890  
891  
892  
893  
894  
895  
896  
897  
898  
899  
900  
901  
902  
903  
904  
905  
906  
907  
908  
909  
910  
911  
912  
913  
914  
915  
916  
917

## 5 Conclusions and Future Work

### 5.1 Conclusions

While more work is needed, deep partitioned neural network architectures using time and frequency data seem promising in long-term solutions for denoising speech and music signals.

### 5.2 Future Work

#### 5.2.1 Models

Make network deeper. Consider gradual partitioning instead of hard.

#### 5.2.2 Data

Get more data. Consider different noise levels and types of signals.